

Contents

1	Fat-pointer Address Translations	1
1.1	Encoding Ranges as Bounds to the Pointer	1
1.2	Instrumenting Block-Based Allocators with Physically Con- tiguous Memory	3
1.3	Sample memory allocator Implementation	3

1 Fat-pointer Address Translations

Fat-pointer Address Translations, combined with the capabilities of the CHERI (Capability Hardware Enhanced RISC Instructions) architecture, introduce robust memory safety and security features by incorporating additional metadata with memory pointers. This enhanced architecture utilizes concepts such as FlexPointer, Range Memory Mapping (RMM) to manage memory effectively.

Range addresses play a pivotal role within this implementation, defining memory regions bounded by a starting address (Upper) and an ending address (Lower). These range addresses are encoded within FAT-pointers, allowing for precise control over memory regions.

Figure 1 illustrates the methodology employed to leverage the CHERI 128-bit FAT-pointer scheme for facilitating block-based memory management on physically contiguous memory, which is depicted on the right side of the figure. This technique contrasts with the conventional approach.

We explore how using Huge pages with CHERI bounds can reduce the number of TLB entries required.

The functionality of ranges encompasses several key aspects:

1.1 Encoding Ranges as Bounds to the Pointer

Integrating range bounds directly into FAT-pointers enables the architecture to enforce memory access restrictions at the pointer level thus allowing tracking of memory ranges on a pointer level. In this implementation, memory ranges are established using bounds encoded within the FAT-pointer, adhering to the CHERI 128-bit bounds compression scheme[1].

Figure 2 illustrates a straightforward use-case in which the dark pink line represents a single, large contiguous memory area, or huge page. Within this huge page, the orange and blue lines indicate two separate memory allocations equivalent to invoking malloc twice to allocate memory in distinct regions. This scenario simulates a block-based memory allocator operating

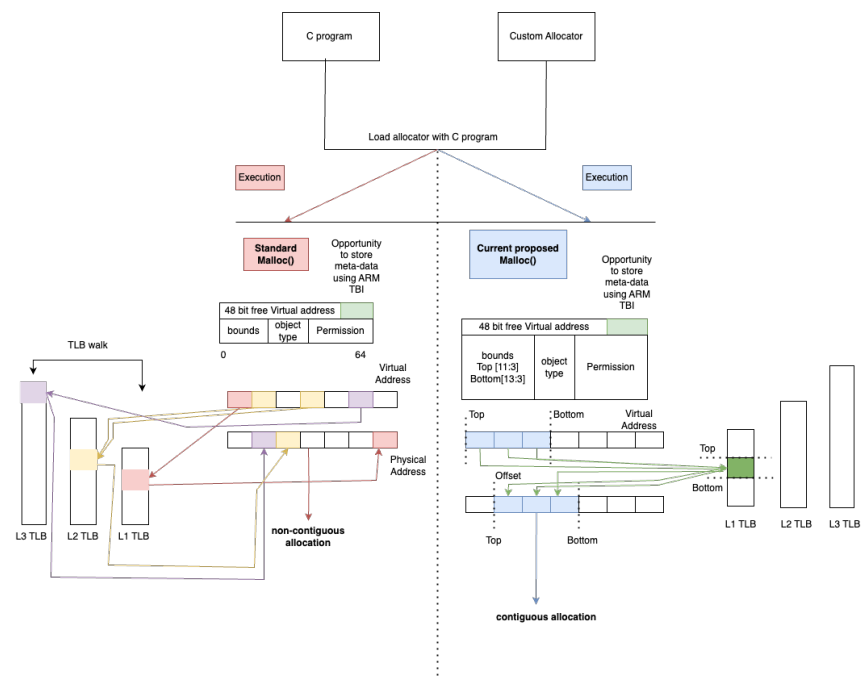


Figure 1: High overview architecture

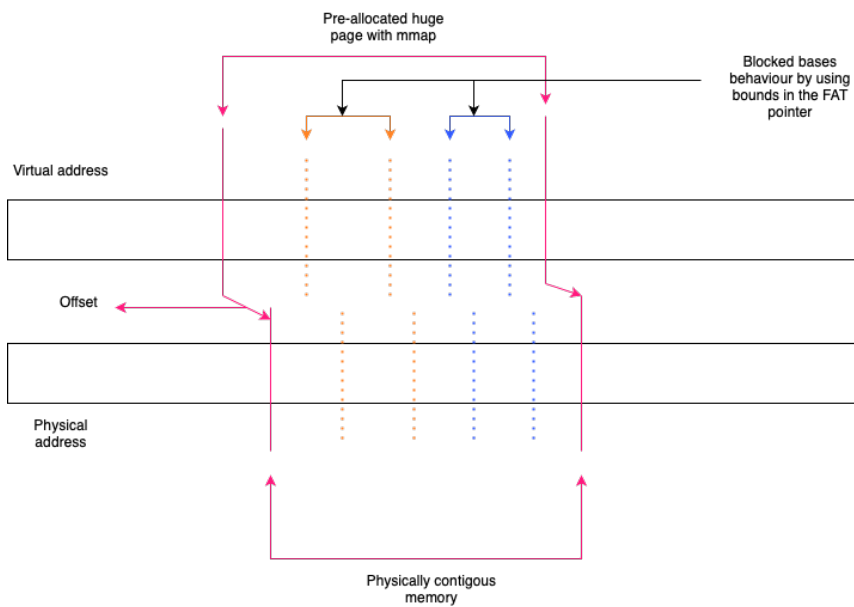


Figure 2: Range of memory

within the confines of the huge page. The allocations leverage the bounds encoded in the FAT-pointer, ensuring tracking and efficient management of the allocated memory regions. By using the FAT-pointer bounds, this method maintains the integrity and contiguity of the allocated blocks within the huge page.

1.2 Instrumenting Block-Based Allocators with Physically Contiguous Memory

Traditional address translation methods rely on hierarchical structures to map virtual addresses to physical addresses. This often requires multiple entries to handle different memory segments, which increases overhead and adds complexity to the translation process. In contrast, the current approach simplifies this by using a single TLB (Translation Lookaside Buffer) entry to translate multiple addresses within a contiguous memory range. This reduces the number of TLB entries needed, making the translation process more efficient and less complex.

By consolidating address translations into a single TLB entry, this method cuts down on the overhead of managing many entries. It also takes advantage of the bounds encoded within fat-pointers to track and access memory more efficiently. This streamlined approach allows for precise and effective memory management, especially within large, contiguous memory regions like huge pages. Overall, it simplifies memory operations while improving performance and reduces TLB overhead by reducing TLB walks.

Figure 3 illustrates a use-case of huge pages where the green line represents a sample access to read within a contiguous space of physical memory. The dotted lines represent the bounds for that particular pointer access. Using bounds stored on the pointer a block based pattern can be replicated on physically contiguous memory.

1.3 Sample memory allocator Implementation

This section presents a straightforward memory allocator designed and implemented based on the principles outlined in our approach. The allocator consists of three core functions: `InitAlloc`, `malloc`, and `free`. The `InitAlloc` function initializes the memory pool, setting up the necessary data structures and metadata required for efficient memory management. The `malloc` function is responsible for allocating a contiguous block of memory of a specified size, while the `free` function deallocates the memory, returning it to the pool for future use.

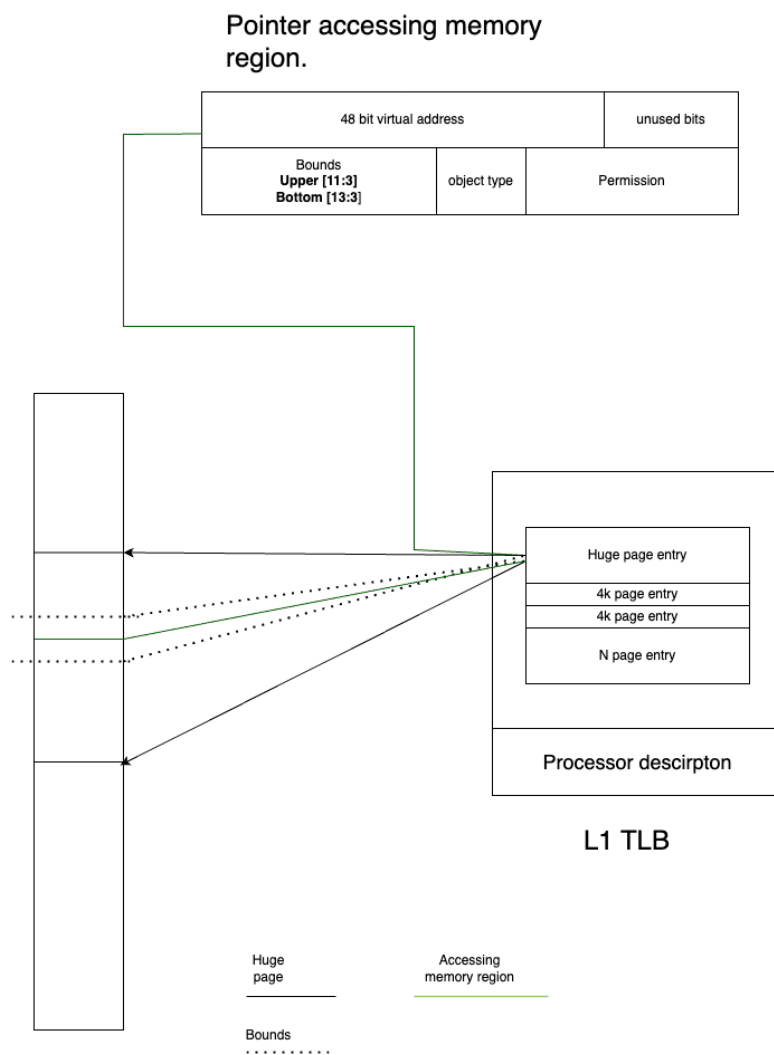


Figure 3: Fat-pointer Address Translations using huge pages

A notable feature of this malloc implementation is its compatibility with kernel modules, where it can be integrated as an alternative to the mmap system call. This integration ensures that memory allocations are physically contiguous, a critical requirement for certain low-level operations and hardware interactions. By providing physically contiguous memory blocks, this allocator can serve as a foundational layer for standard block-based allocators, such as Jemalloc, enabling them to operate efficiently in environments where physical memory contiguity is essential.

Algorithm 1 Sample init alloc function to create a initial 1 GB huge page

```

1: function INIT_ALLOC
2:   sz ← 1 GB                                ▷ Define pre-allocated memory size
3:   fd ← CREATE_LARGE_PAGE_MEMORY(sz)        ▷ Create shared
      memory
4:   ptr ← MAP_MEMORY(sz)                      ▷ Map memory region
5:   MallocCounter ← sz                        ▷ Initialize memory counter
6: end function

```

Algorithm 1 describes the initialization of physically contiguous memory through the use of huge pages, a mechanism supported by modern architectures to optimize memory management. The algorithm begins by allocating a fixed block of 1 GB of physically contiguous memory. This decision is driven by the architectural constraints of contemporary systems, particularly ARM-based CPUs, where 1 GB represents the largest supported page size. By leveraging huge pages, the algorithm reduces the overhead associated with page table management and enhances memory access efficiency, which is critical for performance-sensitive applications and kernel-level operations.

Algorithm 2 Sample malloc implementation

```

1: function MALLOC(sz)
2:   sz ← ALIGN_UP(sz, MAX_ALIGNMENT)        ▷ Align size to max
      alignment
3:   MallocCounter ← MallocCounter - sz      ▷ Update remaining memory
4:   ptrLink ← &ptr[MallocCounter]           ▷ Calculate pointer address
5:   ptrLink ← SET_BOUNDS(ptrLink, sz)       ▷ Set bounds for memory
      safety and to track the length of the pointer
6:   return ptrLink                          ▷ Return allocated memory pointer
7: end function

```

When the malloc function 2 is invoked, the algorithm employs an eager allocation strategy for physical memory. This is achieved through the use of the SetBounds mechanism, which constructs a FAT-pointer a specialized pointer that encodes both the start and end addresses of the allocated memory region within the pointer itself. The start and end addresses correspond to the size of the memory block requested by malloc. This approach introduces a method of memory tracking, where the bounds of the allocated region are explicitly encoded in the address, enabling efficient monitoring and management of memory usage.

Furthermore, this design leverages shared huge page TLB (Translation Lookaside Buffer) entries to map and track memory addresses. By encoding bounds directly into the address, the algorithm ensures that memory accesses remain within the allocated region, thereby enhancing safety and reducing the risk of out-of-bounds errors. This use of FAT-pointers and shared TLB entries not only aligns with the principles of efficient memory management but also demonstrates a practical usecase of huge pages in CHERI.

Algorithm 3 Sample free implementation

```

1: function FREE(ptr)
2:   len ← GET_LENGTH(ptr) ▷ Get length of memory block from the
   defined bounds
3:   UNMAP(ptr, len) ▷ Release memory block
4: end function

```

The memory deallocation 3 mechanism in the proposed allocator is facilitated by the FAT-pointer structure introduced in the malloc algorithm. When the free function is invoked, it utilizes the metadata embedded within the FAT-pointer to determine the range and size of the allocated memory region. Specifically, the start and end addresses encoded in the FAT-pointer provide the necessary information to identify the exact memory block to be deallocated. This allows the allocator to precisely unmapped the corresponding memory region from the address space, ensuring efficient and accurate memory management.

By extracting the bounds and size directly from the FAT-pointer, the free function eliminates the need for additional metadata lookups or complex data structures, streamlining the deallocation process. This approach not only enhances performance but also reduces the risk of memory leaks or fragmentation.

References

- [1] J. Woodruff, A. Joannou, H. Xia, A. Fox, R. M. Norton, D. Chisnall, B. Davis, K. Gudka, N. W. Filardo, A. T. Markettos, M. Roe, P. G. Neumann, R. N. M. Watson, and S. W. Moore, “CHERI concentrate: Practical compressed capabilities,” vol. 68, no. 10, pp. 1455–1469. [Online]. Available: <https://ieeexplore.ieee.org/document/8703061/>